

Problem Set 2

Applied Stats/Quant Methods 1

Due: October 23, 2025

Instructions

- Please show your work! You may lose points by simply writing in the answer. If the problem requires you to execute commands in `R`, please include the code you used to get your answers. Please also include the `.R` file that contains your code. If you are not sure if work needs to be shown for a particular problem, please ask.
- Your homework should be submitted electronically on GitHub.
- This problem set is due before 23:59 on Thursday October 23, 2025. No late assignments will be accepted.

Question 1: Political Science

The following table was created using the data from a study run in a major Latin American city.¹ As part of the experimental treatment in the study, one employee of the research team was chosen to make illegal left turns across traffic to draw the attention of the police officers on shift. Two employee drivers were upper class, two were lower class drivers, and the identity of the driver was randomly assigned per encounter. The researchers were interested in whether officers were more or less likely to solicit a bribe from drivers depending on their class (officers use phrases like, “We can solve this the easy way” to draw a bribe). The table below shows the resulting data.

¹Fried, Lagunes, and Venkataramani (2010). “Corruption and Inequality at the Crossroad: A Multi-method Study of Bribery and Discrimination in Latin America. *Latin American Research Review*. 45 (1): 76-97.

	Not Stopped	Bribe requested	Stopped/given warning
Upper class	14	6	7
Lower class	7	7	1

- (a) Calculate the χ^2 test statistic by hand/manually (even better if you can do "by hand" in R).

Step 1: Put the data in R and calculate the totals

First I put the tabular data into a matrix in R and added the totals row and column. I labeled the matrix the same way the original table was labeled.

```

1 observed_table <- matrix(
2   c(14, 6, 7, 7, 7, 1),
3   nrow = 2,
4   ncol = 3,
5   byrow = TRUE
6 )
7
8 table_row_totals <- rbind(observed_table,
9                            c(sum(observed_table[,1]),
10                             sum(observed_table[,2]),
11                             sum(observed_table[,3])))
12
13 class_table <- cbind(table_row_totals,
14                      c(sum(table_row_totals[1,]),
15                        sum(table_row_totals[2,]),
16                        sum(table_row_totals[3,])))
17
18
19 rownames(class_table) <- c("Upper Class", "Lower Class", "Column Total")
20 colnames(class_table) <- c("Not Stopped", "Bribe Requested",
21                            "Stoped/Given Warning", "Row Total")
22 class_table

```

	Not Stopped	Bribe Requested	Stoped/Given Warning
Upper Class	13.5	8.357143	5.142857
Lower Class	7.5	4.642857	2.857143

Step 2: Calculate the expected values

Find the expected values by calculating row total/grand total * column total

```
1 expected_values <- matrix(c(  
2   (class_table[1,4] / class_table[3,4] * class_table[3,1]),  
3   (class_table[1,4] / class_table[3,4] * class_table[3,2]),  
4   (class_table[1,4] / class_table[3,4] * class_table[3,3]),  
5   (class_table[2,4] / class_table[3,4] * class_table[3,1]),  
6   (class_table[2,4] / class_table[3,4] * class_table[3,2]),  
7   (class_table[2,4] / class_table[3,4] * class_table[3,3])),  
8   nrow = 2,  
9   ncol = 3,  
10  byrow = TRUE  
11 )  
12  
13 rownames(expected_values) <- c("Upper Class", "Lower Class")  
14 colnames(expected_values) <- c("Not Stopped", "Bribe Requested",  
15                               "Stopped/Given Warning")  
16 expected_values
```

	Not Stopped	Bribe Requested	Stopped/Given Warning
Upper Class	13.5	8.357143	5.142857
Lower Class	7.5	4.642857	2.857143

Step 3: Assumptions

- The data is categorical
- The observations are independent
- The groups are mutually exclusive
- All of the expected counts are at least 5 (that isn't met in this example)

Step 4: State hypotheses

- Null Hypothesis: The variables are statistically independent
- Alternative Hypothesis: The variables are statistically dependent

Step 5: Calculate the Chi Squared Statistic

Calculate the Chi-Squared Statistic by taking the (observed values minus the expected values) squared, then dividing that by the expected value. Do that for every value in the table then sum all of them together

```
1 test_statistic <- sum(((observed_table - expected_values)^2)  
2                      / expected_values)  
3 test_statistic
```

```
[1] 3.791168
```

- (b) Now calculate the p-value from the test statistic you just created (in R).² What do you conclude if $\alpha = 0.1$?

Step 6: Calculate the degrees of freedom

Find the degrees of freedom by doing (Rows - 1) * (Columns - 1)

```
1 class_df <- (nrow(observed_table) - 1) * (ncol(observed_table) - 1)
2 class_df
```

```
[1] 2
```

Step 7: Calculate the p-value

Input the test statistic and degrees of freedom in R to get the p-value

```
1 pchisq(test_statistic, df=class_df, lower.tail=FALSE)
```

```
[1] 0.1502306
```

The p-value of 0.1502306 is greater than the critical value of 0.1. Therefore, we fail to reject the null hypothesis that the variables are statistically independent.

²Remember frequency should be > 5 for all cells, but let's calculate the p-value here anyway.

- (c) Calculate the standardized residuals for each cell and put them in the table below.

Step 8: Calculate the standardized residuals

Calculate the standardized residuals by doing the observed minus the expected values divided by the square root of the expected frequency * (1 - the row proportion) * (1 - the column proportion)

Cell [1,1]

```
1 difference <- observed_table - expected_values
2 difference[1,1] / sqrt(expected_values[1,1] *
3   (1 - class_table[1,4] / class_table[3,4]) *
4   (1 - class_table[3,1] / class_table[3,4]))
```

[1] 0.3220306

Cell [1,2]

```
1 difference[1,2] / sqrt(expected_values[1,2] *
2   (1 - class_table[1,4] / class_table[3,4]) *
3   (1 - class_table[3,2] / class_table[3,4]))
```

[1] -1.641957

Cell [1,3]

```
1 difference[1,3] / sqrt(expected_values[1,3] *
2   (1 - class_table[1,4] / class_table[3,4]) *
3   (1 - class_table[3,3] / class_table[3,4]))
```

[1] 1.523026

Cell [2,1]

```
1 difference[2,1] / sqrt(expected_values[2,1] *
2   (1 - class_table[2,4] / class_table[3,4]) *
3   (1 - class_table[3,1] / class_table[3,4]))
```

[1] -0.3220306

Cell [2,2]

```
1 difference[2,2] / sqrt(expected_values[2,2] *  
2   (1 - class_table[2,4] / class_table[3,4]) *  
3   (1 - class_table[3,2] / class_table[3,4]))
```

[1] 1.641957

Cell [2,3]

```
1 difference[2,3] / sqrt(expected_values[2,3] *  
2   (1 - class_table[2,4] / class_table[3,4]) *  
3   (1 - class_table[3,3] / class_table[3,4]))
```

[1] -1.523026

	Not Stopped	Bribe requested	Stopped/given warning
Upper class	0.3220306	-1.641957	1.523026
Lower class	-0.3220306	1.641957	-1.523026

(d) How might the standardized residuals help you interpret the results?

Since none of our standard residuals have an absolute value over 2, we can determine that there are no values significantly outside of what we would expect based on our null hypothesis that the variables are statistically independent. This agrees with the information from our p-value that we cannot reject the null hypothesis that the variables are statistically independent.

Question 2: Economics

Chattopadhyay and Duflo were interested in whether women promote different policies than men.³ Answering this question with observational data is pretty difficult due to potential confounding problems (e.g. the districts that choose female politicians are likely to systematically differ in other aspects too). Hence, they exploit a randomized policy experiment in India, where since the mid-1990s, $\frac{1}{3}$ of village council heads have been randomly reserved for women. A subset of the data from West Bengal can be found at the following link: <https://raw.githubusercontent.com/kosukeimai/qss/master/PREDICTION/women.csv>

Each observation in the data set represents a village and there are two villages associated with one GP (i.e. a level of government is called "GP"). Figure 1 below shows the names and descriptions of the variables in the dataset. The authors hypothesize that female politicians are more likely to support policies female voters want. Researchers found that more women complain about the quality of drinking water than men. You need to estimate the effect of the reservation policy on the number of new or repaired drinking water facilities in the villages.

Figure 1: Names and description of variables from Chattopadhyay and Duflo (2004).

Name	Description
GP	An identifier for the Gram Panchayat (GP)
village	identifier for each village
reserved	binary variable indicating whether the GP was reserved for women leaders or not
female	binary variable indicating whether the GP had a female leader or not
irrigation	variable measuring the number of new or repaired irrigation facilities in the village since the reserve policy started
water	variable measuring the number of new or repaired drinking-water facilities in the village since the reserve policy started

³Chattopadhyay and Duflo. (2004). "Women as Policy Makers: Evidence from a Randomized Policy Experiment in India. *Econometrica*. 72 (5), 1409-1443.

- (a) State a null and alternative (two-tailed) hypothesis.

Step 1: State Hypotheses

1. Null hypothesis: the reservation policy has no affect on the number of drinking water facilities
2. Alternative hypothesis: the reservation policy has an affect on the number of drinking water facilities

- (b) Run a bivariate regression to test this hypothesis in R (include your code!).

Step 2: Run a bivariate regression

Use the `lm` function in R to run the regression, I subset two columns from the original data set to input only the two variables that I am investigating the relationship between, reserved and water.

```
1 women <- read.csv("women.csv")
2 bivariate_regression <- lm(women$water ~ women$reserved)
3 bivariate_regression
```

Call:

```
lm(formula = women$water ~ women$reserved)
```

Coefficients:

```
(Intercept)  women$reserved
14.738      9.252
```

Summary of the regression output

```
1 summary(bivariate_regression)
```

Call:

```
lm(formula = women$water ~ women$reserved)
```

Residuals:

Min	1Q	Median	3Q	Max
-23.991	-14.738	-7.865	2.262	316.009

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	14.738	2.286	6.446	4.22e-10 ***
women\$reserved	9.252	3.948	2.344	0.0197 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 33.45 on 320 degrees of freedom

Multiple R-squared: 0.01688, Adjusted R-squared: 0.0138

F-statistic: 5.493 on 1 and 320 DF, p-value: 0.0197

(c) Interpret the coefficient estimate for reservation policy.

Interpret the coefficients

- A coefficient estimate of 9.252 suggests that for a one unit increase in reservation policy (in this case a one unit increase represents going from no reservation policy to having a reservation policy) there is a 9.252 increase in the number of new or repaired drinking water facilities in the village
- A p-value of 0.0197 is less than the critical value of 0.05, indicating that we can reject the null hypothesis that there is no relationship between reservation policy and the number of drinking water facilities
- A y-intercept of 14.738 suggests that there are an average of 14.738 new or repaired drinking-water facilities when there is no reservation policy in place.